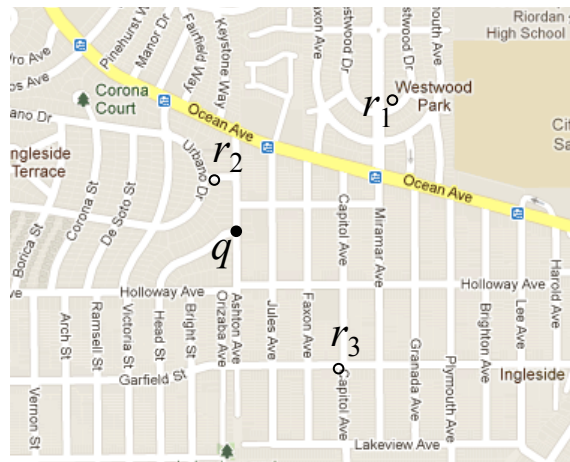


Spatial Networks

- ❑ Modeling and storing spatial networks
- ❑ Shortest path search
- ❑ Spatial queries over spatial networks
- ❑ Advanced indexing techniques for spatial networks




Network Distance

- In many real applications accessibility of objects is restricted by a spatial network

- Examples

- Driver looking for nearest gas station
- Mobile user looking for nearest restaurant




- **Shortest path distance** used instead of Euclidean distance
- $SP(a,b)$ = path between a and b with the minimum accumulated length

Challenges

- ❑ Euclidean distance is no longer relevant
 - R-tree may not be useful, when search is based on shortest path distance
- ❑ Graph cannot be flattened to a one-dimensional space
 - Special storage and indexing techniques for graphs are required
- ❑ Graph properties may vary
 - directed vs. undirected
 - length, time, etc. as edge weights

Modeling Spatial Networks

- Adjacency matrix only appropriate for dense graphs
- Spatial networks are sparse: use adjacency lists instead



graph

	n_1	n_2	n_3	n_4	n_5	n_6
n_1	0	1	∞	∞	∞	∞
n_2	2	0	1	3	4	∞
n_3	∞	∞	0	∞	∞	∞
n_4	∞	∞	∞	0	2	∞
n_5	∞	∞	∞	∞	0	5
n_6	∞	∞	∞	∞	∞	0


adjacency matrix

n_1	$(n_2, 1)$
n_2	$(n_1, 2), (n_3, 1), (n_4, 3), (n_5, 4)$
n_4	$(n_5, 2)$
n_5	$(n_6, 5)$

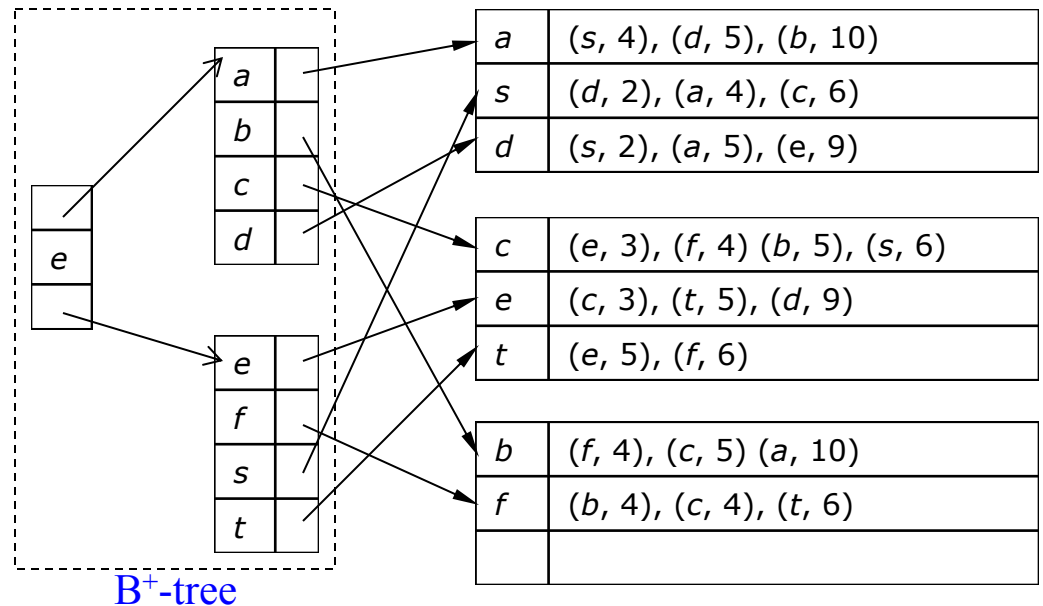
adjacency lists

Storing Large Spatial Networks

- ❑ Problem: adjacency lists representation may not fit in memory if graph is large
- ❑ Solution:
 - partition adjacency lists to disk blocks [based on proximity]
 - create B⁺-tree index on top of partitions [based on node-id]



graph





Shortest Path Computation

- Given a graph $G(V,E)$, and two nodes s,t in V , find the shortest path from s to t
- A classic algorithmic problem
- Studied extensively since the 1950's
- Several methods
 - Dijkstra's algorithm
 - A*-search
 - Bi-directional search
 - *Reach* preprocessing heuristic

Dijkstra's Shortest Path Search

- Idea: incrementally explore the graph around s , visiting nodes in distance order to s until t is found (like NN)





Dijkstra's Shortest Path Search

```
function Dijkstra_SP(source node  $s$ , target node  $t$ )
1.  for each graph node  $v$ 
2.       $SPD(s, v) := \infty$ ; /* initialize shortest path distance */
3.       $path(s, v) := null$ ; /* initialize shortest path */
4.      mark  $v$  as unvisited;
5.  initialize a priority queue  $Q$ ;
6.   $SPD(s, s) := 0$ ; add  $s$  to  $Q$ ;
7.  while not empty( $Q$ )
8.       $v := top(Q)$ ; /* node  $v$  on  $Q$  with smallest  $SPD(s, v)$  */
9.      remove  $v$  from  $Q$ ;
10.     mark  $v$  as visited;
11.     if  $v = t$  then return  $path(s, t)$ ;
12.     for each neighbor  $u$  of  $v$ 
13.         if  $u$  is not marked as visited
14.             if  $SPD(s, u) > SPD(s, v) + weight(v, u)$ 
15.                  $SPD(s, u) := SPD(s, v) + weight(v, u)$ ;
16.                  $path(s, u) := path(s, v) + (v, u)$ ;
17.                 add or update  $u$  on  $Q$ ;
```


Dijkstra's Shortest Path Search

Current node	Queue
s (0)	d (2), a (4), c (6)
d (2)	a (4), c (6), e (11)
a (4)	c (6), e (11), b (14)
c (6)	e (9), f (10), b (11)
e (9)	f (10), b (11), t (14)
f (10)	b (11), t (14)
b (11)	t (14)
t (14)	


t is de-queued: shortest path has been found!



shortest path tree of node s


A*-search

- ❑ Dijkstra's search explores nodes around s without a specific search direction until t is found
- ❑ Idea: improve Dijkstra's algorithm by directing search towards t
- ❑ Due to triangular inequality, Euclidean distance is a lower bound of network distance
- ❑ Use Euclidean distance to lower bound network distance based on known information:
 - Nodes are visited in increasing $\text{SPD}(s,v) + \text{dist}(v,t)$ order
 - ❑ $\text{SPD}(s,v)$: shortest path distance from s to v (computed by Dijkstra)
 - ❑ $\text{dist}(v,t)$: Euclidean distance between v and t
 - Original Dijkstra visits nodes in increasing $\text{SPD}(s,v)$ order



A*-search: Example



Current node	Queue
$s(0)$	$d(2+14), a(4+15), c(6+7)$
$c(6)$	$a(4+15), d(2+14), b(11+9), f(10+6), e(9+5)$
$e(9)$	$a(4+15), d(2+14), b(11+9), f(10+6), t(14)$
$t(14)$	



t is de-queued: shortest path has been found!

Bi-directional search


- ❑ Dijkstra's search explores nodes around s without a specific search direction until t is found
- ❑ Idea: search can be performed concurrently from s and from t (backwards)
- ❑ The shortest path tree of s and the (backward) shortest path tree of t are computed concurrently
 - One queue Q_s for forward and one queue Q_t for backward search
 - Node visits are prioritized based on $\min(\text{SPD}(s,v), \text{SPD}(v,t))$
 - If v already visited from s and v in Q_t , then candidate shortest path: $p(s,v) + (v,u) + p(u,t)$ [if v already visited from t and v in Q_s symmetric] in Q_t
 - If v visited by both s and t terminate search; report best candidate shortest path



Bi-directional search: Example

Current node	Q_s	Q_t
$s(0), t(0)$	$d(2), a(4), c(6)$	$e(5), f(6)$
$d(2) [s]$	$a(4), c(6), e(11)$	$e(5), f(6)$
$a(4) [s]$	$c(6), e(11), b(14)$	$e(5), f(6)$
$e(5) [t]$	$c(6), e(11), b(14)$	$f(6), c(8), d(14)$
$c(6) [s]$	$e(9), f(10), b(11)$	$f(6), c(8), d(14)$
$f(6) [t]$	$e(9), f(10), b(11)$	$c(8), b(10), d(14)$
$c(8) [t]$		

c is visited from both s and t !
 terminate and report shortest path



candidate shortest path:
 $s \rightarrow d \rightarrow e \rightarrow t$ (16)


candidate shortest path:
 $s \rightarrow c \rightarrow e \rightarrow t$ (14)

REACH: A Preprocessing Heuristic

- For every node v in the network find all pairs of nodes s, t such that $SP(s,t)$ includes v
 - Local reach $r_{s,t}(v) = \min\{SPD(s,v), SPD(v,t)\}$
 - **Reach** $r(v) = \max\{r_{s,t}(v)\}$, for all s,t such that $SP(s,t)$ includes v
 - $r(v)$: in the worst case, what is the minimum of $SPD(s,v)$, $SPD(v,t)$ given that v is in $SP(s,t)$?
- Assume $r(v)$ is precomputed for every v in the graph
- In bi-directional search, assume v is visited (e.g. from s)
 - For every unvisited neighbor u of v (by FWD and BK search):
 - $SPD(s,v) \leq SPD(s,u)$ and $SPD(s,v) \leq SPD(u,t)$
 - If $r(u) < SPD(s,v)$ then
 - $r(u) < SPD(s,u)$ and $r(u) < SPD(u,t) \Rightarrow$
 - u cannot be in $SP(s,t) \Rightarrow$
 - u needs not be enheaped (pruned)
- **Reach** reduces the nodes to be added to the heap
 - However, it may require high preprocessing cost (all SPs computed)

Bi-directional w/ Reach: Example

Current node	Q_s	Q_t
$s(0), t(0)$	$c(6)$	$e(5), f(6)$
$e(5) [t]$	$c(6)$	$c(8)$
$c(6) [s]$	$e(9)$	$c(8)$
$c(8) [t]$		



candidate shortest path:
 $s \rightarrow c \rightarrow e \rightarrow t$ (14)


c is visited from both s and t !
 terminate and report shortest path

Combination of Techniques

- ❑ A*, bi-directional search, and **Reach** can be combined to a powerful search technique
- ❑ A* can only be applied if lower distance bounds are available
- ❑ **Reach** can only be applied after pre-processing
 - expensive
 - high maintenance cost
- ❑ All versions of Dijkstra's search require non-negative edge weights
 - Bellman-Ford is an algorithm for arbitrary edge weights (some of them could be negative)

Source/Destination on Edges

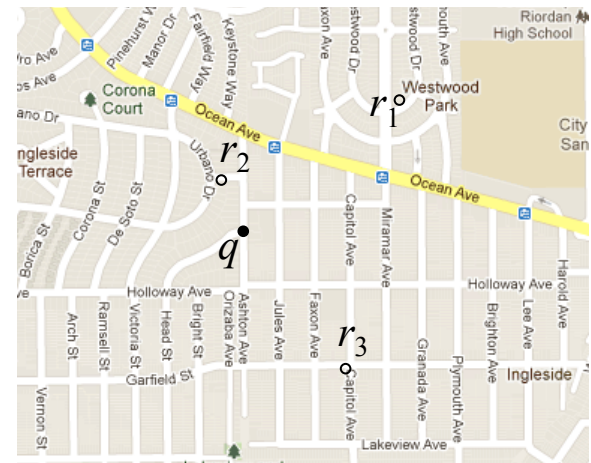

- We have assumed that points s and t are nodes of the network
- In practice s and t could be arbitrary points on edges
 - Mobile user locations
- Solve problem by introducing 2 more nodes



- en-heap g (2) and c (4) first
- t is reached from f or h

Spatial Queries over Spatial Networks

- Data:
 - A (static) spatial network (e.g., city map)
 - A (dynamic) set of spatial objects
- Spatial queries based on network distance:
 - Selections. Ex: find gas stations within 10km driving distance from here
 - Nearest neighbor search. Ex: find k nearest restaurants from present position
 - Joins. Ex: find pairs of restaurants and hotels at most 100m from each other



Methodology



- ❑ Store (and index) the spatial network
 - Graph component (indexes connectivity information)
 - Spatial component (indexes coordinates of nodes, edges, etc.)
- ❑ Store (and index) the sets of spatial objects
 - Ex., one spatial relation for restaurants, one spatial relation for hotels, one relation for mobile users, etc.
- ❑ Given a spatial location p , use spatial component of network to find the network edge containing p
- ❑ Given a network edge, use network component to traverse neighboring edges
- ❑ Given a neighboring edge, use spatial indexes to find objects on them

Evaluation of Spatial Selections (1)

- ❑ Query: find all objects in spatial relation R , within network distance ε from location q
- ❑ Method:
 - Use spatial index of network (R-tree indexing network edges) to find edge n_1n_2 , which includes q
 - Use adjacency index of network (graph component) and apply Dijkstra's algorithm to progressively retrieve edges that are within network distance ε from location q
 - For all these edges apply a spatial selection on the R-tree that indexes R to find the results



Evaluation of Spatial Selections (1)

- Example: Find restaurants at most distance 10 from q
- Step 1: find network edge which contains q





Evaluation of Spatial Selections (1)

- Example: Find restaurants at most distance 10 from q
- Step 2: traverse network to find all edges (or parts of them) within distance 10 from q



Evaluation of Spatial Selections (1)

- Example: Find restaurants at most distance 10 from q
- Step 3: find restaurants that intersect the subnetwork computed at step 2





Evaluation of Spatial Selections (2)

- ❑ Query: find all objects in spatial relation R , within network distance ε from location q
- ❑ Alternative method based on Euclidean bounds:
 - Assumption: Euclidean distance is a lower-bound of network distance:
 - ❑ $\text{dist}(v,u) \leq \text{SPD}(v,u)$, for any v,u
 - Use R-tree on R to find set S of objects such that for each o in S : $\text{dist}(q,o) \leq \varepsilon$
 - For each o in S :
 - ❑ find where o is located in the network (use Network R-tree)
 - ❑ compute $\text{SPD}(q,o)$ (e.g. use A^*)
 - ❑ If $\text{SPD}(q,o) \leq \varepsilon$ then output o



Evaluation of Spatial Selection (2)

- Example: Find restaurants at most distance 10 from q
- Step 1: find restaurants for which the Euclidean distance to q is at most 10: $S = \{r_1, r_2, r_3\}$



Evaluation of Spatial Selections (2)

- Example: Find restaurants at most distance 10 from q
- Step 2: for each restaurant in S, compute SPD to q and verify if it is indeed a correct result
 - Ex. for r_2 , first find where r_2 is located (edge f,t), then SP





Evaluation of NN search (1)

- ❑ Query: find in spatial relation R the nearest object to a given location q
- ❑ Method:
 - Use spatial index of network (R-tree indexing network edges) to find edge n_1n_2 , which includes q
 - Use adjacency index of network (graph component) and apply Dijkstra's algorithm to progressively retrieve edges in order of their distance to q
 - For each edge apply a spatial selection on the R-tree that indexes R to find any objects
 - Keep track of nearest object found so far; use its shortest path distance to terminate network browsing

Evaluation of Spatial Selections (1)

- Example: Find nearest restaurant to q




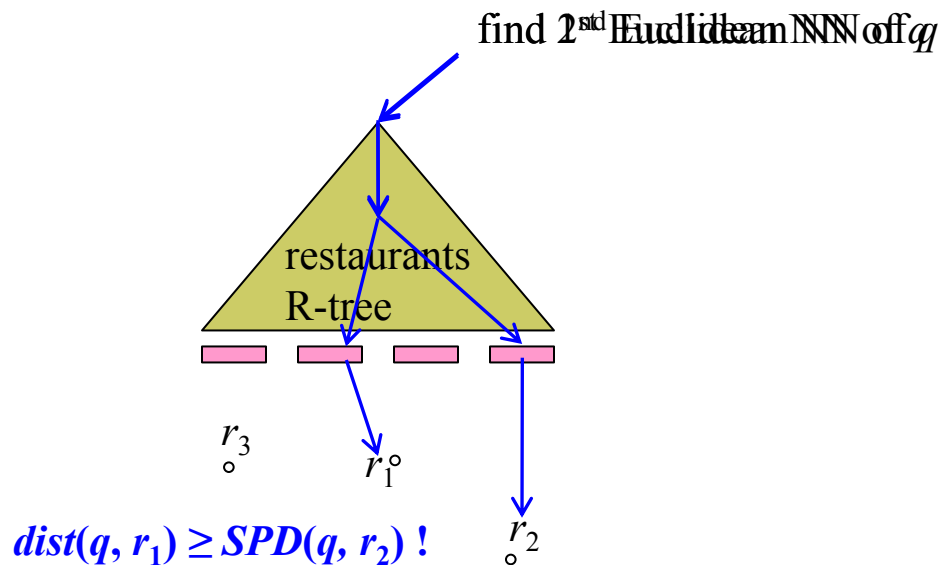
current NN = r_2 ✓
 $SPD(q, r_2) = 5$

Evaluation of NN search (2)

- ❑ Query: find in spatial relation R the nearest object to a given location q
- ❑ Alternative method based on Euclidean bounds:
 - Assumption: Euclidean distance lower-bounds network distance:
 - ❑ $\text{dist}(v,u) \leq \text{SPD}(v,u)$, for any v,u
 - 1 Use R-tree on R to find Euclidean NN p_{E1} of q;
 - 2 CurrentNN= p_{E1} ; bound=SPD(q, p_{E1}); //e.g. use A*
 - 3 Find next Euclidean NN p_{Ei} of q
 - 4 If $\text{dist}(q,E_i) \geq \text{bound}$, then
 report (CurrentNN,bound) as result;
 - 5 Compute SPD(q, p_{Ei}); if $\text{bound} > \text{SPD}(q,p_{Ei})$ then
 CurrentNN= p_{Ei} ; bound=SPD(q, p_{Ei});
 - 6 Goto step 3

Evaluation of Spatial Selections (2)

- Example: Find nearest restaurant to q



current NN = r_2 ✓
 $SPD(q, r_2) = 5$

Spatial Join Queries


- ❑ Query: find pairs (r,s) , such that r in relation R , s in relation S , and $SPD(r,s) \leq \epsilon$
- ❑ Methods:
 - For each r in R , do an ϵ -distance selection queries for objects in S (Index Nested Loops)
 - For each pair (r,s) , such that $\text{Euclidean dist}(r,s) \leq \epsilon$ compute $SPD(r,s)$ and verify $SPD(r,s) \leq \epsilon$

Notes on Query Evaluation based on Network Distance

- For each query type, there are methods based on network browsing and methods based on Euclidean bounds
- Network browsing methods are fast if network edges are densely populated with points of interest
 - A limited network traversal can find the result fast
- Methods based on Euclidean bounds are good if the searched POIs are sparsely distributed in the network
 - Few verifications with exact SP searches are required
 - Directed SP search (e.g. using A^*) avoids visiting empty parts of the network

Shortest Path Materialization and Indexing in Large Graphs

- ❑ Dijkstra's algorithm and related methods could be very expensive on very large graphs
- ❑ (Partial) materialization of shortest paths in static graphs can accelerate search



graph

	<i>a</i>	<i>b</i>	<i>c</i>	<i>s</i>	...
<i>a</i>	0, -	10, <i>ab</i>	10, <i>asc</i>	4, <i>as</i>	...
<i>b</i>	10, <i>ba</i>	0, -	5, <i>bc</i>	11, <i>bcs</i>	...
<i>c</i>	10, <i>csa</i>	5, <i>cb</i>	0, -	6, <i>cs</i>	...
<i>s</i>	4, <i>sa</i>	11, <i>scb</i>	6, <i>sc</i>	0, -	...
...

brute-force materialization
 $O(n^3)$ space, $O(1)$ time


	<i>a</i>	<i>b</i>	<i>c</i>	<i>s</i>	...
<i>a</i>	0, -	10, <i>b</i>	10, <i>s</i>	4, <i>s</i>	...
<i>b</i>	10, <i>a</i>	0, -	5, <i>c</i>	11, <i>c</i>	...
<i>c</i>	10, <i>s</i>	5, <i>b</i>	0, -	6, <i>s</i>	...
<i>s</i>	4, <i>a</i>	11, <i>c</i>	6, <i>c</i>	0, -	...
...

distance matrix with successors
 $O(n^2)$ space, $O(n)$ time

Hierarchical Path Materialization

- ❑ Idea: Partition graph G into G_1, G_2, G_3, \dots based on connectivity and proximity of nodes
- ❑ Every edge of G goes to exactly one G_i
- ❑ **Border nodes** belong to more than one G_i 's
- ❑ For each G_i compute and materialize SPs between every pair of nodes in G_i (matrix M_i)
 - Partitions are small enough for materialization space overhead to be low
- ❑ Compute and materialize SPs between every pair of border nodes (matrix B)
 - If border nodes too many, hierarchically partition them into 2nd-level partitions

Hierarchical Path Materialization




border nodes SP materialization


Hierarchical Path Materialization

- ❑ Shortest path search using HPM:
- ❑ If s, t border nodes directly use matrix B
- ❑ If s border, t non-border
 - $SP(s, t) = \min\{p(s, u) + p(u, t) \mid u \in B \text{ and } u \in G_t\}$
- ❑ If s non-border, t border
 - $SP(s, t) = \min\{p(s, u) + p(u, t) \mid u \in B \text{ and } u \in G_s\}$
- ❑ If s, t non-border nodes
 - If s, t in same G_{st} then $SP(s, t)$ materialized in M_{st}
 - If s in G_s , t in G_t then
 - ❑ $SP(s, t) = \min\{p(s, u) + p(u, v) + p(v, t) \mid u, v \in B, u \in G_s, v \in G_t\}$


Hierarchical Path Materialization



$SP(s,t) = ?$




□ *border nodes*
● *internal nodes*



border nodes SP materialization

Hierarchical Path Materialization

- ❑ Good partitioning if:
 - small partitions
 - few combinations examined for SP search
- ❑ Real road networks:
 - Non-highway nodes in local partitions
 - Highway nodes become border nodes



Compressing Materialized Paths



- Distance matrix with successors has $O(n^2)$ space cost
- Motivation: reduce space by grouping targets based on common successors

	<i>a</i>	<i>b</i>	<i>c</i>	<i>s</i>	...
<i>a</i>	0, -	10, <i>b</i>	10, <i>s</i>	4, <i>s</i>	...
<i>b</i>	10, <i>a</i>	0, -	5, <i>c</i>	11, <i>c</i>	...
<i>c</i>	10, <i>s</i>	5, <i>b</i>	0, -	6, <i>s</i>	...
<i>s</i>	4, <i>a</i>	11, <i>c</i>	6, <i>c</i>	0, -	...
...

distance matrix with successors
 $O(n^2)$ space, $O(n)$ time



Compressing Materialized Paths

- ❑ Create and encode one space partitioning defined by targets of the same successor
- ❑ For each node s , index I_s a set of $\langle \text{succ}, R \rangle$ pairs:
 - succ : a successor of s
 - R : a continuous region, such that for each t in R , the successor of s in $\text{SP}(s,t)$ is succ




Compressing Materialized Paths

- To compute $SP(s,t)$ for a given s, t :
 1. $SP = s$
 2. Use spatial index I_s to find $\langle succ, R \rangle$, such that t in R
 3. $SP = SP + (s, succ)$
 4. If $succ = t$, report SP and terminate
 5. Otherwise $s = succ$; Goto step 2



Embedding Methods for SP

- ❑ Select a small subset L of nodes in G
- ❑ Nodes in L are called **landmarks**
- ❑ For each l in L , for each v in G :
 - pre-compute $\text{SPD}(v, l)$, $\text{SPD}(l, v)$
- ❑ To compute $\text{SP}(s, t)$ use A^* search
 - To compute the lower bound distance between any node v and t we can use **(triangular inequality)**:
 - ❑ $D = \max\{|\text{SPD}(v, l) - \text{SPD}(l, t)|, \text{ for each } l \text{ in } L\}$
 - Prioritization of node visits is now done using $\text{SPD}(s, v) + D$
- ❑ Ex: if r_3 is used as landmark, then $|\text{SPD}(q, r_3) - \text{SPD}(r_3, r_2)|$ is a better LB for $\text{SPD}(q, r_2)$ compared to Euclidean distance



Summary

- ❑ Indexing and search of spatial networks is different than spatial indexing
 - Shortest path distance is used instead of Euclidean distance, to define range queries, nearest neighbor search, and spatial joins
- ❑ Spatial networks could be too large to fit in memory
 - Disk-based index for adjacency lists is used
- ❑ Several shortest path algorithms
- ❑ Spatial queries can be evaluated using Euclidean bounds
- ❑ Advanced indexing methods for shortest path search on large graphs